

CS 311O

Mutable Data Types

A New Despair
Mutability Strikes Back
Return of Imperative Programming

Nate Foster

Spring 2020

Today's music: *The Imperial March*
from the soundtrack to *Star Wars, Episode V: The Empire Strikes Back*

CLICKER QUESTIONS 1 AND 2

Review

Previously in 3110:


- Efficiency: Big Oh

Lectures 13-15:

- Efficiency of data structures
- Running example: maps

Today: THE DARK SIDE ARRIVES

- Mutable data types: refs, mutable fields, (arrays)



REFS

Demo

References


- Aka “ref” or “ref cell”
- **Pointer** to a typed location in memory
- Binding of variable to pointer: **immutable**
- Contents of memory location: **mutable**

References

- Syntax: **ref e**
- Evaluation:
 - Evaluate **e** to a value **v**
 - Allocate a new *location* **loc** in memory to hold **v**
 - Store **v** in **loc**
 - Return **loc**
- Type checking:
 - New type constructor: **t ref** where **t** is a type
 - Note: **ref** is used as keyword in type and as keyword in value
 - **ref e : t ref** if **e : t**

Locations

- Locations are values
- Locations are not expressions
- So this picture [lec02] was a lie:



References

- Syntax: $e1 := e2$
- Evaluation:
 - Evaluate $e2$ to a value $v2$
 - Evaluate $e1$ to a location loc
 - Store $v2$ in loc
 - Return $()$
- Type checking:
 - If $e2 : t$
 - and $e1 : t \text{ ref}$
 - then $e1 := e2 : \text{unit}$

Unit

- `unit` is a type
 - Its only value is `()`, also pronounced “unit”
 - There are no interesting operations on `unit`
- Analogy: Booleans
 - `bool` is a type
 - It has just two values, `true` and `false`
 - So `unit` is like a Boolean with one fewer value
- Analogy: `void`
 - When a procedure in Java or C has no interesting value to return, its return type is `void`
 - Similar with `print` and `assert` in OCaml:
 - `print_string : string -> unit`
 - `assert b : unit` (assuming `b : bool`)

References

- **Syntax: $!e$**
 - note: not negation
- **Evaluation:**
 - Evaluate e to loc
 - Return contents of loc
- **Type checking:**
 - If $e : t \text{ ref}$
 - then $!e : t$

CLICKER QUESTION 3

Equality

- Suppose we have two refs...
 - `let r1 = ref 3110`
 - `let r2 = ref 3110`
- Double equals is *physical equality*
 - `r1 == r1`
 - `r1 != r2`
- Single equals is *structural equality*
 - `r1 = r1`
 - `r1 = r2`
 - `ref 3110 <> ref 2110`
- You usually want single equals

Semicolon

- Syntax: **`e1 ; e2`**
- Evaluation:
 - Evaluate **`e1`** to a value **`v1`**
 - Then throw away that value
(note: **`e1`** could have side effects)
 - evaluate **`e2`** to a value **`v2`**
 - return **`v2`**
- Type checking:
 - If **`e1 : unit`**
 - and **`e2 : t`**
 - then **`e1 ; e2 : t`**

MUTABLE FIELDS

Implementing refs

Ref cells are essentially implemented as records with a mutable field:

```
type 'a ref = { mutable contents : 'a }  
let ref x = { contents = x }  
let (!) r = r.contents  
let (:=) r newval = r.contents <- newval
```

- That type is declared in **Stdlib**
- The functions are compiled down to something equivalent



BEWARE


Immutable lists

We have never needed to worry about aliasing with lists!


```
let x = [2;4]
```

```
let y = [5;3;0]
```

```
let z = x @ y
```



vs.



(no code you write could ever tell, but OCaml implementation uses the first one)

OCaml:

blissfully unaware of aliasing

Java:

obsession with aliasing

Faulty code

```
class ProtectedResource {
    private Resource theResource = ...;
    private String[] allowedUsers = ...;
    public String[] getAllowedUsers() {
        return allowedUsers;
    }
    public String currentUser() { ... }
    public void useTheResource() {
        for(int i=0; i < allowedUsers.length; i++) {
            if(currentUser().equals(allowedUsers[i])) {
                ... // access allowed: use it
                return;
            }
        }
        throw new IllegalAccessException();
    }
}
```

Can you find the security fault?

Have to make copies

The exploit:

```
p.getAllowedUsers()[0] = p.currentUser();  
p.useTheResource();
```

The fix:

```
public String[] getAllowedUsers() {  
    ... return a copy of allowedUsers ...  
}
```

Similar errors as recent as Java 1.7beta

Pros and cons of immutability

Pros:

- Programmer doesn't have to think about aliasing; can concentrate on other aspects of code
- Language implementation is free to use aliasing, which is cheap
- Often easier to reason about whether code is correct

Cons:

- I/O is fundamentally about mutation
- Some data abstractions (dictionaries, arrays, ...) are more efficient if imperative

Try not to abuse your new-found power!

Upcoming events

- [last night] R4 due
- [Wed] A3 due
- Check CodePost for pinned exam-related posts!

This is (reluctantly) imperative.

THIS IS 3110