



# CS 3110

## Balanced Trees

Nate Foster  
Spring 2020

Today's scene: Libe Slope

# Review

## Previously in 3110:

- Efficiency
- Hash tables: imperative constant-time maps

## Today:

- Balanced trees: functional maps

# Running example: Sets

```
module type Set = sig
  type 'a t
  val empty : 'a t
  val insert : 'a -> 'a t -> 'a t
  val mem : 'a -> 'a t -> bool
  ...
end
```


**LIST SET**

Demo

**BST SET**

# Binary search tree (BST)

- Binary tree: every node has two subtrees
- BST invariant:
  - all values in  $l$  are less than  $v$
  - all values in  $r$  are greater than  $v$



# Set implementations: performance

|         | Workload 1 |      |
|---------|------------|------|
|         | insert     | mem  |
| ListSet | 35s        | 106s |

MacBook, 1.3 GHz Intel Core m7, 8 GB RAM, median of three runs

# Set implementations: performance

|         | Workload 1 |      |
|---------|------------|------|
|         | insert     | mem  |
| ListSet | 35s        | 106s |
| BstSet  | 130s       | 149s |

MacBook, 1.3 GHz Intel Core m7, 8 GB RAM, median of three runs



# Set implementations: performance

|         | Workload 1 |      | Workload 2 |       |
|---------|------------|------|------------|-------|
|         | insert     | mem  | insert     | mem   |
| ListSet | 35s        | 106s | 35s        | 106s  |
| BstSet  | 130s       | 149s | 0.07s      | 0.07s |


MacBook, 1.3 GHz Intel Core m7, 8 GB RAM, median of three runs

# Workloads

- Workload 1:
  - insert: 50,000 elements in ascending order
  - mem: 100,000 elements, half of which not in set
- Workload 2:
  - insert: 50,000 elements in random order
  - mem: 100,000 elements, half of which not in set


# Insert in random order

- Resulting tree depends on exact order
- One possibility for inserting 1..4 in random order  
3, 2, 4, 1:



# Insert in linear order

Only one possibility for inserting 1..4 in linear order  
1, 2, 3, 4:




unbalanced: leaning toward the right

# When trees get big

- Inserting next element in linear tree **always** takes  $n$  operations where  $n$  is number of elements in tree already
- Inserting next element in randomly-built tree **might** take far fewer...

# Best case tree



all paths through *perfect binary tree* have same length:  $\log_2 (n+1)$ ,  
where  $n$  is the number of nodes,  
recalling there are implicitly leafs below each node at bottom level

# Performance of BST

- `insert` and `mem` are both  $O(n)$
- But if trees always had short paths instead of long paths, could be better:  $O(\log n)$
- How could we ensure short paths?  
i.e., *balance* trees so they don't lean




# **BALANCED TREES**



# Strategies for achieving balance

- In general:
  - Strengthen the RI to require balance
  - And modify insert to guarantee it
- Well-known data structures:
  - 2-3 trees: all paths have **same length**
  - AVL trees: length of shortest and longest path from any node **differ at most by one**
  - Red-black trees: length of shortest and longest path from any node **differ at most by factor of two**
- All of these achieve  $O(\log n)$  insert and mem



# CS 3110

## RED-BLACK TREES


# Red-black trees

- [Guibas and Sedgwick 1978], [Okasaki 1998]
- Binary search tree with:
  - Each node colored red or black
  - Leaves and root colored black
- RI: BST +
  - **Local invariant:** No red node has a red child
  - **Global invariant:** Every path from the root to a leaf has the same number of black nodes

# Path length

- Invariants:
  - No red node has a red child
  - Every path from the root to a leaf has the same number of black nodes
- Together imply: length of longest path is **at most twice** length of shortest path
  - e.g., B-**R**-B-**R**-B-**R**-B vs. B-B-B-B

# Examples of new nodes







New red node  
violates Local  
Invariant

New black  
node violates  
Global  
Invariant


# Okasaki's algorithm

- [Okasaki 1998]: functional RB tree
- Always maintain BST + Global Invariant
- Make new node red
- Recurse back up tree
  - Look at the two nodes immediately beneath current node
  - Fix any violations of Local Invariant with a **rotation** that balances tree

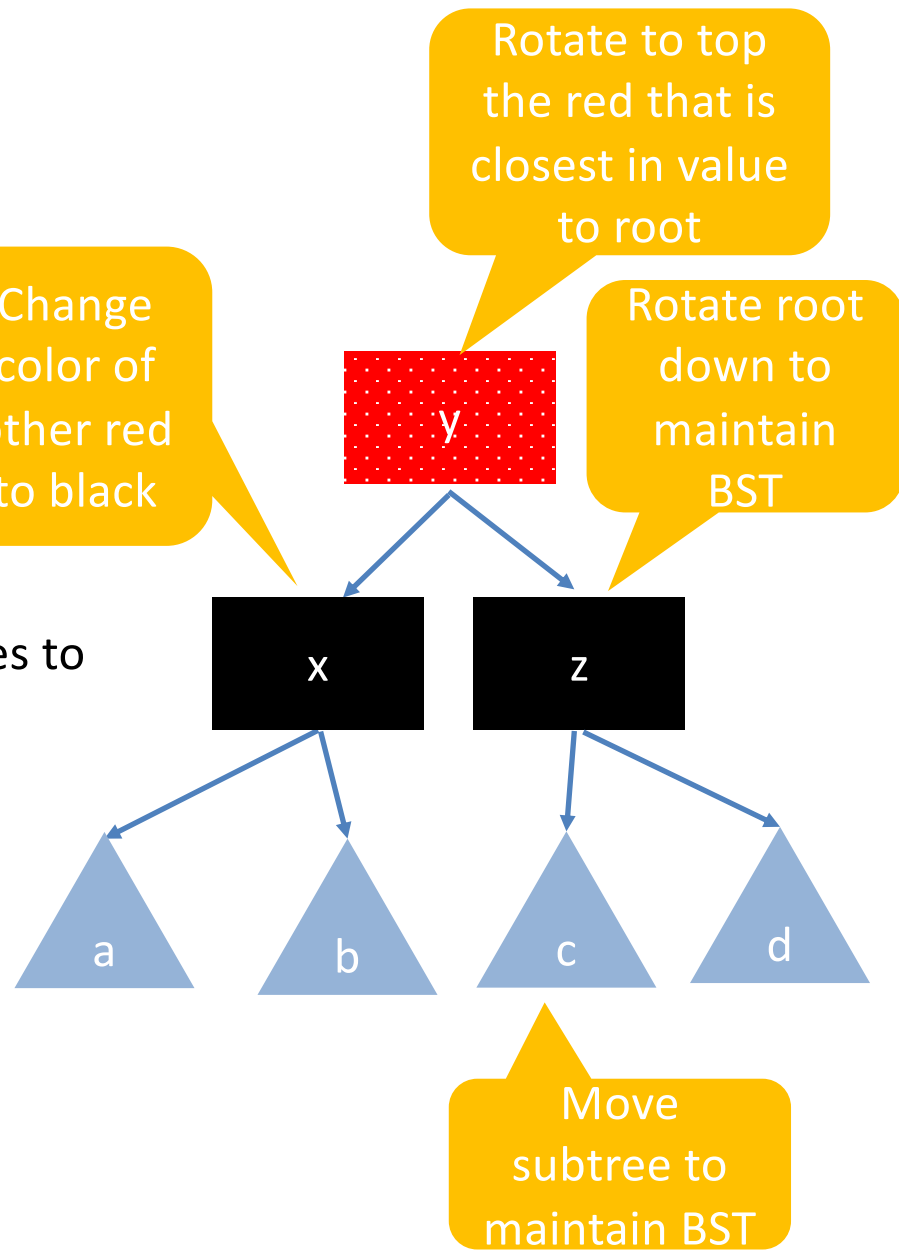
# Only four possible violations



# RB rotate (1 of 4)




rotates to






# RB rotate (1 of 4)



rotates to




BST + Global Invariant: maintained

Local Invariant: fixed for x & y


...but maybe broken for y and z's original parent!

...so keep recursing up


# RB rotate (2 of 4)




rotates to




# RB rotate (3 of 4)




rotates to



# RB rotate (4 of 4)



rotates to



# OCaml implementation

```
let balance = function  
  | (Blk, Node (Red, Node (Red, a, x, b), y, c), z, d) (* 1 *)  
  | (Blk, Node (Red, a, x, Node (Red, b, y, c)), z, d) (* 2 *)  
  | (Blk, a, x, Node (Red, Node (Red, b, y, c), z, d)) (* 4 *)  
  | (Blk, a, x, Node (Red, b, y, Node (Red, c, z, d))) (* 3 *)  
    -> Node (Red, Node (Blk, a, x, b), y, Node (Blk, c, z, d))  
  | t -> Node t
```




# Efficiency of red-black set

- mem:  $O(\log n)$ 
  - Worst case: walk down one entire path
  - Length of every path is the same as a perfect tree, or at most twice that long
- insert:  $O(\log n)$ 
  - You will see algorithm in recitation (and textbook)
  - Worst case: walk down then up one entire path

# Red-black dictionary

- Store (key, value) pair at each node
- Order as a BST by keys



# Map implementations

|                   | insert      | find        | remove      |
|-------------------|-------------|-------------|-------------|
| Arrays            | $O(1)$      | $O(1)$      | $O(1)$      |
| Association lists | $O(1)$      | $O(n)$      | $O(n)$      |
| Hash tables       | $O(1)$      | $O(1)$      | $O(1)$      |
| Red-black trees   | $O(\log n)$ | $O(\log n)$ | $O(\log n)$ |

- **Arrays:** fast, but keys must be integers
- **Association lists:** allow any keys, but slower
- **Hash tables:** fast, but requires good hash function; and worst-case performance relaxed to expected & amortized performance
- **Red-black trees:** almost as fast, and immutable




# Set implementations: performance

|         | Workload 1 |       | Workload 2 |       |
|---------|------------|-------|------------|-------|
|         | insert     | mem   | insert     | mem   |
| ListSet | 35s        | 106s  | 35s        | 106s  |
| BstSet  | 130s       | 149s  | 0.07s      | 0.07s |
| RbSet   | 0.12s      | 0.07s | 0.15s      | 0.08s |

MacBook, 1.3 GHz Intel Core m7, 8 GB RAM, median of three runs

# Sir Tony Hoare



b. 1934

Turing Award Winner 1980

*For his fundamental contributions to the definition and design of programming languages.*

"We should forget about small efficiencies, say about 97% of the time: **premature** optimization is the root of all evil."

# Upcoming events

- [Monday] A4 released
- [Tuesday] Discussion sections start
- [Wednesday] Form partners on CMS
- [Friday] MS0 due

*This is blissfully balanced.*

**THIS IS 3110**